


HW3 - Report

Q1

The analytical partial derivative w.r.t b :

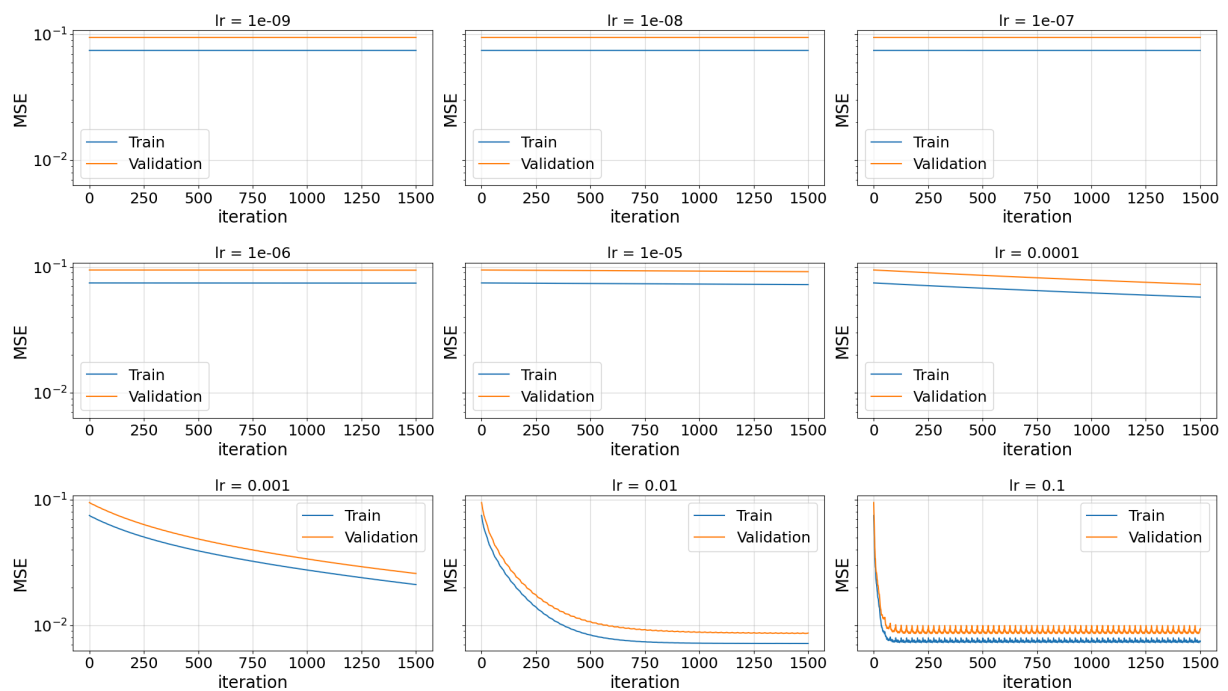
$$\frac{\partial}{\partial b} L(w, b) = \frac{1}{m} 2 \cdot (1_m)^T (Xw + 1_m \cdot b - y)$$

Q2



Q3

Training and Validation Losses as a Function of Iteration # for Different LRs



The optimal learning rate is **0.01** because its training and validation loss decreased the most and steadily compared to other learning rates.

Learning rates $1e-9$ to $1e-5$ do not demonstrate a decrease in the losses. This is because these learning rates are very small, making the updates to the model parameters insignificant. Thus the model does not learn and improve at all over the 1500 iterations.

For learning rates 0.0001 we see a minimal decrease in the losses, which suggests that the model does learn a little but the learning rate is still relatively too low for significant progress. This is also true for learning rate 0.001 for which the losses decrease even more but are still much worse than the learning rate 0.01.

For the learning rate 0.1, its losses initially decreased significantly but then they ceased to improve, and the loss jumps a lot which indicates that 0.1 is too large as a learning rate, causing the model to make very large updates to its parameters instead of converging to a local minimum.

We see that 0.01 learning rate is enough to reach minimum loss even in less than 1500 steps, and increasing the number of steps will not improve the model.

Q4


Crossed-validated errors of regressors

Model	Section	Train MSE	Valid MSE
		Cross validated	
Dummy	2	0.068899	0.068933

Q5

learning rates tested values: `np.logspace(-6, -1, 10)`

Number of folds in CV: 5



The optimal learning rate is **0.0278**

Crossed-validated errors of regressors

Model	Section	Train MSE	Valid MSE
		Cross validated	
Dummy	2	0.068899	0.068933
Linear	2	0.007412	0.007722


Q6

The performance of the Dummy Regressor would not change, because it always returns the mean of the target value, which is not affected by feature normalization. However, the linear regressor performance would decrease, because it relies on minimizing the MSE, and if some features are on a way larger scale, they will contribute a lot to the loss, and the regressor would fine-tune to minimize them, rather than taking into consideration all the given features equally.

Q7

Regularization strength values: `np.logspace(-4, 0, 10)`

Number of folds in CV: 5



The optimal regularization strength: **0.000774**

Crossed-validated errors of regressors


Model	Section	Train MSE	Valid MSE
		Cross validated	
Dummy	2	0.068899	0.068933
Linear	2	0.007412	0.007722
Lasso Linear	3	0.0074	0.007686

Q9

Top 5 features with the largest coefficients (in absolute value) in the optimal Lasso regressor

Feature	Coefficient
PCR_04	0.462614
sugar_levels	0.126135
PCR_02	-0.077948
PCR_06	-0.047372
PCR_09	-0.010335

Q10



Q11

The magnitude of the coefficients is interesting because it indicates the feature importance¹, i.e. its influence strength on the target variable. Larger magnitudes (in absolute value) mean a greater impact. By comparing feature coefficients we can determine what features contribute to the model more than others, aiding in feature selection. In particular, a zero coefficient indicates that the feature has no contribution to the model and can be considered for removal.

Here we see that the model could still reach high accuracy by just using the top 5 features as stated in Q9, as the rest have very low coefficients.

Q12


The training performance of the Lasso model would have changed. Similarly to the Linear Regressor, when the features aren't normalized, differently scaled features can dominate the loss expression, and the model would have to fine-tune them to decrease the loss, resulting in missed patterns that could be found by using all the features.

Q13

I would expect the coefficients of a Ridge regressor to be more uniformly small (but not almost 0). This is because the Lasso regressor uses L1 regularization, which encourages sparse weight vector (and we can see in Q10 that indeed most features have almost 0 coefficient). This can be said to induce “feature selection”. On the contrary, the Ridge regressor uses L2 regularization, which penalizes the sum of the squared coefficients. This encourages the overall shrinkage of all coefficients towards zero uniformly (since we normalized the features).

Section 4 Task 2:

Visualization of contamination_level vs PCR_04 and PCR_06



We can see it looks pretty linear (sits on a plane), except when looking at the right-most perspective where we can see some quadratic behavior.

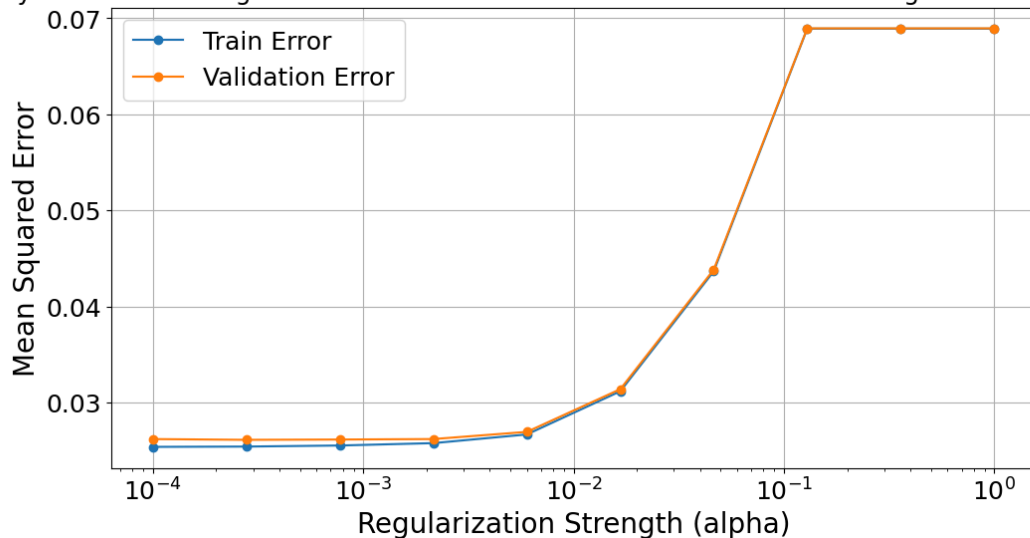
¹ This works because the features are normalized. Had the features not been normalized, the coefficients would be biased because features could be in different scales, as discussed in Q6.

Q14

It is important to re-normalize because if we have a monomial x , and we compare it to x^2 , for big values of x , x^2 will be a lot bigger. If we don't normalize, the higher-order monomials will have way bigger values than the rest and will dominate the loss expression just like if we don't normalize the features from the input. To use higher-order monomials we normalize them, which lets them keep their behavior (curvature, shape, and so on) while letting all monomials be of equal magnitude, which would be better for the linear regressor to train with.

Q15

Polynomial-Lasso Regressor Train and Validation Errors as a function of Regularization Strength




Optimal regularization strength: **0.000278**

Optimal validation error: 0.026

Q16

Visualization of contamination_level vs PCR_04 and PCR_06



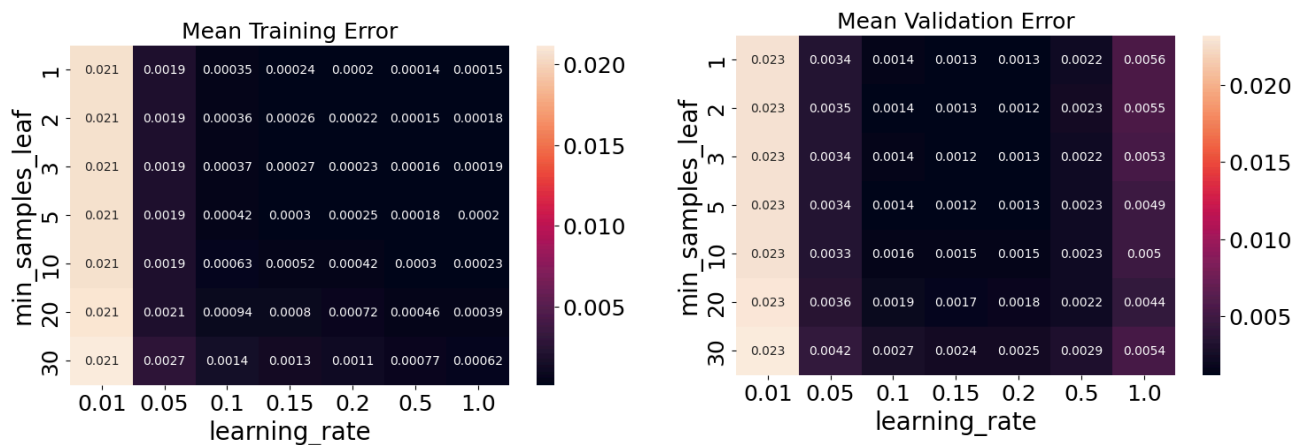
The points in blue are the model's predictions and the points in red are the train-dataset. We see the model fits the data pretty well, which makes sense since the data appears to lie on a plane.

Q17

Crossed-validated errors of regressors

Model	Section	Train MSE	Valid MSE
		Cross validated	
Dummy	2	0.068899	0.068933
Linear	2	0.007412	0.007722
Lasso Linear	3	0.0074	0.007686
Polynomial Lasso	4	0.025422	0.026124

Q18



Optimal parameters: **Learning Rate 0.15, Min Samples Leaves: 3**

Optimal parameters train error: 0.00027477, validation error: 0.001192

Q19

Crossed-validated errors of regressors

Model	Section	Train MSE	Valid MSE
		Cross validated	
Dummy	2	0.068899	0.068933
Linear	2	0.007412	0.007722
Lasso Linear	3	0.0074	0.007686
Polynomial Lasso	4	0.025422	0.026124
GBM Regressor	5	0.000275	0.001192

Q20**Crossed-validated and Test errors of all regressors**

Model	Section	Train MSE	Valid MSE	Test MSE
		Cross validated		Retrained
Dummy	2	0.068899	0.068933	0.065379
Linear	2	0.007412	0.007722	0.007742
Lasso Linear	3	0.0074	0.007686	0.007673
Polynomial Lasso	4	0.025422	0.026124	0.023838
GBM Regressor	5	0.000275	0.001192	0.001502

The **GBM** model performed best on the test set.

We can see that in terms of overfitting/underfitting:

The GBM model was overfitting (lower train error than test error by an order of 10) but still had the lowest test error of all models. Other models had similar train/valid/test errors. Among them, the Polynomial Lasso can be considered the most underfitting (except the dummy regressor of course) since it has the highest errors, which are also closest to the dummy regressor. This suggests that the polynomial mapping did not contribute to the Lasso model performance. It is also interesting to note that the Linear and the Lasso Linear models achieved very similar MSE in the Train/Valid/Test sets.